

# The RNN API

Sequences present several complexities.

Let's begin by better understanding functions that have sequences as inputs and output.

We call this the [RNN API \(RNN\\_API.ipynb\)](#).

# Inside an RNN layer

By now you hopefully have a good intuitive understanding of a Recurrent Layer, but lack the details

Let's open up the hood and [go inside an RNN \(RNN Workings.ipynb\)](#).

# RNN in action

A concrete example may help you to appreciate the power of an RNN.

The task we solve is called Sentiment Analysis

- Given a sequence of words
- Is the sentiment express Positive or Negative ?

In particular

- The examples are movie reviews from IMdB

## IMdb examples

This notebook is from the (future) module on NLP.

The beginning of this notebook addresses the aspects particular to dealing with Natural Language

- [the input examples \(Keras examples imdb cnn.ipynb#Examine-the-text-data\)](#)
- the pre-processing steps necessary for NLP:
  - breaking character strings into words
  - mapping words to integers (index within a finite vocabulary)

To summarize our approach to dealing with words

- We have a finite vocabulary  $\mathbf{V}$
- Words are One Hot Encoded
- So the input sequence  $\mathbf{x}_{(1)} \dots \mathbf{x}_{(T)}$  is
  - a sequence of length  $T$
  - of OHE vectors of length  $\|\mathbf{V}\|$

We also have the option of transforming the OHE word encodings into *Embeddings*

- Embeddings are a transformation of the OHE into shorter vectors that *also* encode some "meaning" to words
- Our initial pass: we just use an Identity transformation
- This will be covered in detail in the NLP module.

We will cover the NLP aspects in more detail in the NLP module.


For now: we will focus on using an RNN to create a finite length representation of the (potentially unbounded length) sequence of words.

The model we create follows the paradigm in the introduction

- Using an RNN to create a fixed length encoding of a variable length sequence
- A Head Layer that is a Binary Classifier

---

RNN Many to one; followed by classifier



Let's look at the code. The main take-away

- the RNN is a layer just like any other
- we use it as the first layer in a Sequential model
  - to reduce the sequence to a finite vector
- once we have a finite vector, Dense layers further transform the representation
- until we have a final binary Classifier "head"

[RNN\(LSTM\) model \(Keras examples imdb cnn.ipynb#LSTM-w/o-One-Hot-Encoding-the-input:-what-happens-?\)](#)

# What is *really* going on inside an RNN

At this point

- You appreciate the ability of an RNN to operate on sequences
- Understand the mechanics of the internal workings

But the update equations don't really convey an intuition about *how* the RNN achieves its power.



Let's try to visualize the latent state of an RNN in order to get a better grasp.

[RNN Visualization \(RNN\\_Visualization.ipynb\)](#)

# **RNN practicalities**

## **Sequences: Variable length**

There are lots of small potholes one encounters with sequences.

What if the examples of my training set have widely varying lengths ?

- Within a batch, short examples may behave differently than long examples:
  - Maybe learn less in short examples, noisier gradient updates
- Padding sequences to make them equal length
  - Pad at the start ? Or at the end ?

The general advice is to arrange your data so that an epoch contains examples of similar lengths.

- You may require multiple fittings, one per length

# Long sequences

We will learn that long sequences present a challenge to training RNN's

- vanishing gradients
- back propagation of gradients takes a long time

There is also the practical matter of long sequences (e.g., greater than the "max" length allocated to a variable).

A Deeper Dive deals with the practical treatment of [long sequences](#)  
([RNN Long Sequences.ipynb](#)).

# Issues with RNN's

Although an RNN layer seems powerful (and a little magical) we have glossed over some big issues

- Can they handle *long* sequences or are they subject to "forgetting" ?
  - Short term versus long term memory trade offs
- Can we really unroll a computation over a long sequence ?
  - Gradient computation potentially more difficult in very deep graphs
- What are the practical difficulties in Keras with long sequences

These will be the topics of subsequent modules.

- Some topics require an in-depth understanding of Gradient Computation (still to come !)

# Conclusion

The Recurrent layer was yet another layer type that we have introduced in rapid succession.

We chose to do this as a "sprint" rather than a "marathon" so that you can start coding and experimenting.

Use the opportunity ! This is where the real learning will happen.

Our next topics will be a more in-depth exploration of issues that may not have come into view during the sprint.

In [2]: `print("Done")`

Done



